



Governed product analytics for credit-card teams

An LLM-powered copilot that turns business questions into reviewable evidence—without letting the model calculate the numbers.

LLM-guided • Python-verified • Human-approved

The bottleneck is translation—not arithmetic

Product questions lose time and clarity before analysis even begins.

01 Business intent is ambiguous

A PM asks where customers are stuck—but does not specify a metric, funnel, or test.

02 Data meaning is fragmented

Equivalent lifecycle events may use different field names and definitions across teams.

03 Answers require governed evidence

Results must preserve calculations, assumptions, limitations, and human decision authority.

FinSight bridges business language and approved analytics—without becoming a general-purpose agent.

Three roles share one reviewable decision flow

Each user enters with a different job—but leaves with the same governed evidence.

PRODUCT MANAGER

Turn a broad product question into a measurable decision

Reviewable first-pass answer + next step

PRODUCT ANALYST

Standardize recurring funnel and experiment work

Reusable calculations + audit trail


GROWTH / ONBOARDING LEAD

Find friction and evaluate a redesigned journey

Measure → Diagnose → Experiment → Decide

Four use cases cover the credit-card lifecycle

The dataset activates only the lifecycle modules supported by confirmed fields.



DATASET COMPATIBILITY

Ready runs the full workflow • Limited shows what is missing • Unavailable refuses unsupported analysis

Six governed analyses answer recurring product questions

Every workflow returns deterministic outputs before the LLM writes the explanation.

KPI DEFINITION

Primary and supporting metrics

Definitions + current values

FUNNEL

Where customers drop off

Counts + conversion + loss

SEGMENTATION

Device, channel, customer segment

Comparable counts + rates

ENGAGEMENT & SPEND

Usage frequency and depth

Active rate + transactions + spend

RETENTION & INACTIVITY

30/90-day activity and reactivation

Separate window rates

A/B EXPERIMENT

Did the redesign improve activation?

Lift + CI + p-value + SRM + MDE

Trust comes from separating reasoning and calculation

The model never becomes the source of numerical truth.



Audit trail preserves source metadata, confirmed mappings, analysis plan, and deterministic result

A pilot should prove speed and trust

FinSight prepares decisions without promising autonomous analysis or rollout authority.

FASTER

Question → reviewed
first-pass insight

CONSISTENT

Metric and experiment
definitions

SAFER

Fewer unsupported
fields and claims

AUDITABLE

Visible plan, result, and
limitations

Pilot proof points

Routing accuracy • Numerical agreement • Review time • Correction rate • Analyst acceptance

The product thesis: a narrower analytics copilot can create more trust than an open-ended agent.